

# Dojo Field Cards — six habits that fit in a pocket

CSPS AI Training Workshop · 24–28 August 2026 · cut along the borders if you like

## 1 · The AI reads only what is on the desk 机上

What it sees when it answers: your message, what you pasted, and any tool result — nothing else. The whole conversation is re-sent every turn; it remembers nothing between turns.

- New task → new chat (or restate the essentials).
- Paste the document **first**, put the instruction **last**.
- Long texts can be read; complex judgements should be asked in short, well-placed instructions.

After 『MCP完全入門』(2026) and 蒲生 『コンテキストエンジニアリング』(2026).

## 2 · Why long jobs fail 九分

$$0.9 \times 0.9 \times \dots (10 \text{ steps}) \approx 0.35$$

If each step is right 90% of the time, ten steps in a row come out right about one time in three. At 99% per step the chain reaches 90%.

- Cut the job into parts: Role · Task · Context · Constraints · Format.
- Put your own check where the chain is weakest.
- “AI drafts, you decide” is arithmetic, not modesty.

After S. Douglas in *The Scaling Era* (Patel, 2025).

## 3 · Five parts is a ceiling, not a floor 上限

A simple request needs only the deliverable. Add a part when the answer disappoints — one at a time, then compare.

- Avoid the five wasteful prompts: “make it nice” · ten orders at once · missing background · only “don’t...” · using the result unchecked.
- Keep the winners in a prompt library: the prompt / the situation / a ◎○△ mark.
- Your own best document, pasted as the example, beats any clever wording.

After 佐藤 『Claude仕事術』(2026), 池田 『Gemini最強のAI仕事術』(2025).

## 4 · Before you upload: four kinds of file 三分

| File                     | Rule                                      |
|--------------------------|-------------------------------------------|
| Your own material        | Upload freely, within its data class      |
| Purchased / copyrighted  | Private use only; never a shared notebook |
| NDA or “no cloud” clause | Check the contract first                  |
| Contains names or IDs    | Strip identifiers, then upload            |

One confidential line makes the whole prompt confidential. “Not used for training” is a vendor setting, not your organisation’s permission.

After 増田 『NotebookLM即効活用大全』(2026), 上田 · 小坂 『AIエージェン ト... トータルガイド』(2026).

## 5 · The panic phrase & the 60-point landing 一言

*“One moment — let me check my notes.”*

Keep a five-line sticky note, not a script. When the AI’s suggested wording is not something you could actually say:

*“Two shorter, politer versions, please.”*

The target is what *you* can deliver, not 100 points. For a vague request from above: “I take this to mean X — I’ll proceed that way unless you say otherwise.”

After チームシャイニー編 『生成AI仕事術』(翔泳社 2026) — universal design, for everyone.

## 6 · Find one thing to fix 正見

Before any output leaves your desk:

- Read it looking for **one** error. There is usually one.
- Reach the key fact by a second route — the source, a calculation, a colleague. Fluent text that looks like a source is not a source.
- For references: does it exist · where was it accepted · who cites it.
- Then sign: *AI-assisted draft, verified by [name]*.

After 菊田 『原論文から解き明かす生成AI』(2025) and the dojo’s own Art of True Sight.

## Remember: content is data, not orders

Anything the AI reads — an email, a PDF, a web page — can carry hidden instructions (“ignore the task above and...”). Never connect an assistant to your inbox or files without that rule in place; never paste a password or login link into a chat; a human presses **send**, **delete**, **publish** and **pay**. After 『[入門] LLMアプリ開発』(2026) ch.4; 大澤 『claude code実用入門』(2026).