Output one — thinking it through

Policy ideas and research questions: the AI as a sparring partner, never as an oracle.

The failure mode nobody warns you about

It agrees with you. Ask "is this a good research question?" and you will get encouragement, in fluent academic English, for almost anything. That is not judgement — it is the machine predicting what a helpful reply looks like. An assistant that never disagrees is worse than no assistant, because it feels like validation.

The fix: ask for preconditions, not approval

```
#role a sceptical reviewer in my field
#task list the assumptions under which this claim holds
#constraints order them by how easily each could fail; name the one most likely to break first
#format numbered list, one line each
```

Then the follow-ups that actually move a project: *"Is the comparison fair?"* · *"What would have to be true for the opposite conclusion?"* · *"What evidence would change your answer?"*

Three lenses, three different arguments

Run the same idea past three roles in three separate chats. The disagreement between them is where your paper's contribution usually hides.

```
#role a sceptical peer reviewer — where is this thin?
#role a finance ministry official — what would you refuse to fund, and why?
#role a citizen affected by this policy — what does this miss about my life?
```

Diverge before you converge

Ask for **ten framings of the question**, not one answer. Reject nine. The ten cost you two minutes; the rejection is your expertise, and it is the part that cannot be delegated.

✅ The check: **which objection had you not already thought of?** If none, the session produced nothing — sharpen the prompt, or you were talking to a mirror.

Small-country discipline in brainstorming

The moment you describe a real case in enough detail to be useful, you may have identified the people in it. Brainstorm at the level of **the mechanism**, not the case: "a district where one officer handles all applications", not the district's name. Nothing confidential enters a public tool, including in a question.

What stays yours: choosing the question, and deciding which objection matters. The machine can list ten; only you know which one your institution will actually be asked about.

Output two — papers, journals, policy briefs

AI accelerates every step of the writing loop. Authorship never moves.

The loop, in order

```
1 REVERSE-OUTLINE — "list the claims in this paragraph in order, then tell me which sentence carries the main claim."
2 REWRITE — "rewrite so the main claim leads and the evidence follows. Do not add content."
3 SELF-CRITIQUE — "act as a sceptical peer reviewer: your three strongest objections."
4 YOUR FINAL PASS — always. The last edit is the author's.
```

Step 2's constraint is the important one. "Do not add content" is what stops a polished paragraph from acquiring a claim you never made and cannot defend.

Citations: the one place never to delegate

It can invent references complete with plausible authors, journals and page numbers. Three checks before any citation enters a draft:

1. **Does it exist?** Search the exact title. A fabricated entry is the lesson, not a scandal.
2. **Where was it accepted?** For a preprint: published since, and where?
3. **Who cites it?** Ask for the three hub papers and why, not twenty titles.

The [CHECK] convention

```
#constraints where you are not certain of a field, write [CHECK] instead of guessing
```

Uncertainty made visible is uncertainty you can resolve. An annotated bibliography with eleven [CHECK] marks is far more useful than a confident one with three silent errors.

A brief is not a short paper

A paper argues a **contribution**; a brief supports a **decision**. Ask for the brief explicitly, or you will get a compressed paper that no minister can act on.

```
#role policy analyst writing for a decision-maker with four minutes
#task turn this paper into a one-page brief
#constraints lead with the decision at stake; three findings, each with its source page; name one option considered and rejected; no methods section
#format one page, headed sections, no citations in the body
```

Disclosure, in 2026

Journals differ, but three points are common ground: **AI cannot be an author**; use must be **disclosed**; the human authors remain **fully responsible**. Two things are new this year: *Nature* recommends submitting the **prompt-and-response log** the way you submit a dataset where AI touched the creative work; and text from new models now carries an **invisible watermark** (EU AI Act). A detection is a hint, not proof — which is exactly why you declare it yourself, first, in your own words.

✅ The check before submission: every number traced to its source page, every citation opened at least once, and the disclosure sentence written before the deadline panic.

What stays yours: the claim, and the signature under it. Everything above is drafting; none of it is authorship.

Page 3: turning your own research into teaching materials.

Output three — lectures, seminars, handouts

The fastest return of the three: your own finished work, taught.

From your paper to five slides

```
#task turn the attached paper into a five-slide teaching outline
#constraints each title states the takeaway, not the topic; 3 bullets maximum per slide; one speaker note each; cite the source page
#format a table I can paste into PowerPoint
```

Set the expectation: what comes back is an outline and notes — pasting it into PowerPoint is your two minutes. That is the real workflow, and it is still an afternoon saved.

The rule that makes slides teach

One message per slide, and the title states it. Not "Requests by district" but "Three districts absorb four-fifths of the workload". If the title is a topic, the audience must do the work of concluding — and they will each conclude something different.

Build the study materials from your own handouts

Upload your lecture notes to a grounded notebook and use the **Studio** panel: it generates **quizzes** and **flashcards** from your material only, a **recap audio** for the drive home, an **infographic** for the noticeboard, and a **slide deck** draft. Every morning quiz in this workshop was made exactly that way — nothing here was beyond you.

```
Five questions from these notes, for a student who attended the lecture but took no notes.
Ten flashcards, term on one side, one-sentence definition on the other.
A ten-minute discussion between two voices, for students to listen to before the seminar.
```

✅ The check: **verify every number the generated material prints.** Grounded means traceable, not true — if your own notes carry an error, the quiz will teach it to thirty students at once.

Two rules for teaching about your own country

- **Never let it invent a national statistic for a lecture.** Ask for the indicator, then fetch the value yourself from the primary source, and put source and year on the slide. A fabricated figure repeated by thirty students is a real cost, and it will be attributed to you.
- **Label every illustrative case as illustrative.** "Fictional example, built for this class" costs four words and protects both the people in the real case and your standing.

What stays yours: deciding what the students must be able to *do* afterwards. No tool can choose that, and every good slide follows from it.